

Spatial queries – matching candidates


- Very slow for larger tables: $M \times N$ comparisons, each possibly proportional to the number of vertices of participating geometries;
- Retrieval from **hard disk** is much slower than from **memory**.
- We need a **strategy to reduce search space to reduce search time**.
- FILTER → REFINE:**
 - Filter:** Coarse, approximate search – on disk
 - Refine:** Detailed matching and checking for the relationship – in memory if possible (this may include checking all vertices in a geometry).
- Such strategies are applied also in non-spatial contexts

Indexes – retrieval complexity

- Records with close values should be close together in an index (close index values);
- In **clustered indices**, they are also **stored nearby** (speed up!)
- In tables, direct lookup (sequential scan) is of linear complexity (proportional to the number of records, $O(n)$);
- Indexing allows to reduce this to non-linear time (e.g., $O(\log(n))$).
- Order in which you specify indices on your composite index matters – search for `city|city_part|suburb` OR `suburb|city_part|city` ?

Tree data structures

Trees (cont.): terminology



- Index of Left child: $2i+1$
 - Index of Right child: $2i+2$
 - Index of parent: $(i-1)/2$
- [i is index of THIS node]

Tree traversal

- ☒ Depth first search
- ☐ Breadth first search

Indexing

- Index:** a data structure that improves the speed of retrieval of data matching a query;
- Speeds up the identification of records** based on a given **attribute** – to **not** inspect **all** records for **exact** value;
- Aim:** sub-linear time lookup $<O(n)$
- Imposes an **ordering on records** based on the values of an **attribute**:
 - this ordering can be **logical** only, or
 - Clustered index:** if ordering reflected in storage as **physical** ordering (**only one per table**).
- An index stores pointers to the page with the data item**
 - Key:value store of pageid.itemid, **sorted by whichever attribute of the tuple** (with some nuances, more to come)
- Cons:**
 - Indices have to be maintained (re-built);
 - Indices take additional storage space;
 - In some cases, the index match may be very inaccurate (see later)

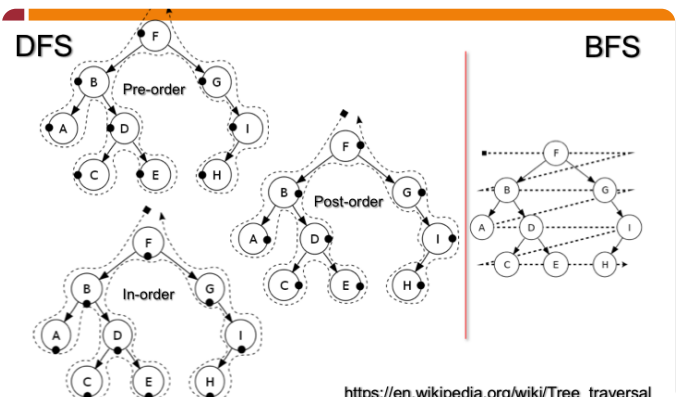
Index: sacrifice storage space, in favour of execution time.

Usage:

- To enforce database constraints (e.g., UNIQUE, PRIMARY KEY);
- To improve lookup on any other column frequently used in WHERE, ORDER BY, and JOIN clauses
- Multi-column indexes are also possible

Trees (cont.): Properties

- Full tree:** every node has 2 children
- Complete tree:** all levels filled except possibly bottom row, filled from left to right



Trees: B⁺-Tree

- B-Trees: family of trees - generalisation of the binary search tree enabling for more children (e.g., B⁺Tree, B*Tree, BSP tree...)
- Order of tree: num children = $m \leq M$ (M often called b)
- Children can be leafs or internal nodes
- Root has at least one key, non-root nodes at least $M/2$ subtrees
- All leaves are sorted, map to disk space (page)
- All leaves are at the same level

Node Type	Children Type	Min Number of Children	Max Number of Children	Example $b = 7$	Example $b = 100$
Root Node (when it is the only node in the tree)	Records	1	$b - 1$	1-6	1-99
Root Node	Internal Nodes or Leaf Nodes	2	b	2-7	2-100
Internal Node	Internal Nodes or Leaf Nodes	$\lceil b/2 \rceil$	b	4-7	50-100
Leaf Node	Records	$\lceil b/2 \rceil$	b	4-7	50-100

In the following, assume

- ▶ M branching factor (max number of children per node, also often b)
- ▶ Consider tree of order 4 ($M=4$)
- ▶ Root node should only have one entry
- ▶ Every non-leaf node should have at least m children nodes: $M/2 < m$ entries per node in leaves $\leq M$
- ▶ $M-1$ search key values in internal nodes
- ▶ Will be a balanced tree
- ▶ The leaf pages must store enough records to remain at least half full

Trees: B⁺-Tree

- B Tree: Keys and records both can be stored in the internal as well as leaf nodes.
- B+ tree: extension of the idea. Records (data) can only be stored on the leaf nodes, while internal nodes can only store the key values.

B-Tree Properties

- ◆ A B-tree is completely balanced (path from root to leaf is constant) at all stages in its evolution
- ◆ Search time is bounded by the length of the path, and so is $O(\log n)$
- ◆ Insertion and deletion of records require $O(\log n)$ time
- ◆ Each node is guaranteed to be at least half full (or almost half full with odd fan-out ratios) at all stages in a B-tree's evolution

B+-trees, where pointers to records are only stored at leaf nodes, are more often used in practice

Spatial indexes

- Generalisation of the previous ideas into 2+ dimensions
- We need to sort things to order and number them in space somehow.
- We want close things to be close together in the index;
- Not natural way to sort x and y coordinate!

Spatial indexes - common approach

- We divide space or group objects based on some criteria;
- These spatial divisions or groups are somehow numbered (indexed);
- The numbering is meaningful, in order to support sequential search;
- Problem:
 - large objects, long objects, concave objects
- Solution: hierarchical divisions

Two approaches:

- Space-driven indexing: space is regularly partitioned into regions (cells) with an assigned index value).
- Data (content) driven indexing
- Problems:
 - We have points, lines, irregular polygons
 - Sizes and spatial distribution of objects can vary hugely
 - Application type: frequent writes or frequent reads?
 - Recent development: Is the DB partitioned?

Compromises are necessary, know your priorities

- What kind of searches (range, exact,...)
- How often we insert new objects;
- How costly is the update of the index.

Space-driven spatial indexes

- We divide space into regular cells;
- These cells may be hierarchically organised;
- Cells may be always present (index slots always stored), or they may be generated based on the data inserted – but the scheme pre-defines the division.


Grid index

- We divide space into regular cells;
- Objects get and address in this space – either numbered, or as x and y index
- One object may be assigned to multiple cells
- Sorting is based on binary sorts (e.g., B-Tree)
- Numbering of cells following some logical order:
 - Row ordering (below, aka row-prime)
- Pros:
 - Simple application of 1D sorting;
 - Fast insert
- Cons:
 - Inefficient for storage size, may include large amounts of empty, pre-allocated space or redundancy
 - Depending on object distribution less efficient range search (object close in space may be far in index);

Object driven indexes

- Adapt to the distribution of objects (approximated by bounding shapes) in the plane/space;
- Do not fully cover space – only space occupied by spatial objects;
- Based on containment relationship, recursively
- Rely on a balanced hierarchical structure with direct mapping into a disk space (page)

Quadtrees: variants (Region, Point – KD tree)



Quad tree index

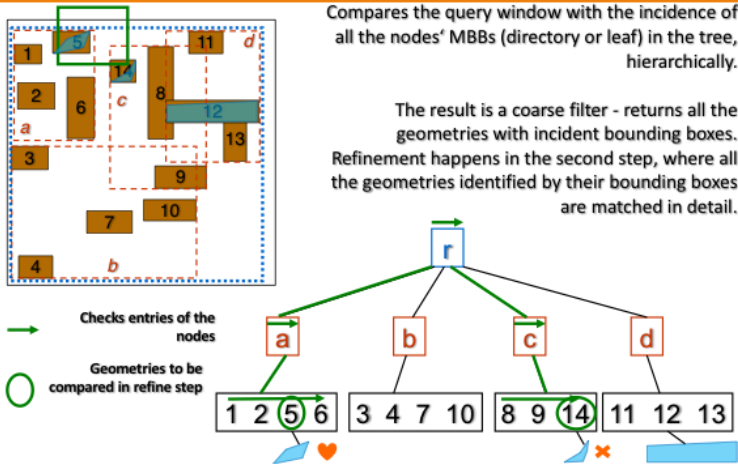
- Divide space into hierarchically nested, adaptable cells;
- Each cell has a limit size (N of objects)
- The directory (the index array) follows a spatially motivated numbering scheme.
- Recursive (self-similar) numbering – space filling curves
- Pros:
 - Simple to index using modifications of standard 1D sorting;
 - Useful for raster indexing!
 - Adapt better to data (less wasted space)
 - Requires only local recomputing after insert/deletion
- Cons:
 - Depending on object distribution, may be less efficient for range search (object close in space may still be far in index);
 - Trees can be unbalanced in depth – hard to estimate retrieval time;
 - Some quadtrees depend on the order in which objects are inserted

R-Tree indexes

- Various refinements: R*Tree, R+Tree, STR-Tree...
- Main principle:
 - Depth-balanced tree (maps well to B+Tree) – all leaves on the same level;
 - Each node corresponds to a disk page
 - Each leaf node contains an array of leaf entries [(mbb: o_id)];
 - Non-leaf nodes contain an array of node entries [(mbb: node_id)];
 - Number of entries in node is $m < x < M$, where m in $[0, M/2]$


Filter-Refine data lookup

R-Tree indexes – Lookup/spatial search




Rectangles and Minimum Bounding Boxes

- Minimum bounding box (MBB/MBR): the smallest rectangle bounding a shape with its axes parallel to the sides of the Cartesian frame
- Using MBB, some queries may be answered without retrieving the geometry of an object
- E.g., find all objects which lie entirely within a specified region




R-Tree indexes: Problem cases

R-Trees can lead, after many I/Os to some problem cases, deteriorating the quality of the tree:



R-Tree indexes: Problem cases

Solution: recompute tree in its entirety:

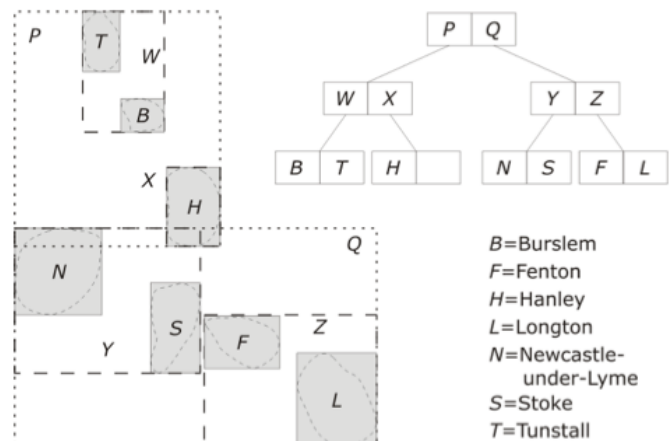


These cases are the focus of the R*Tree, R+Tree and the SRT Tree

R-Tree

- Multidimensional dynamic spatial data structure similar to the B-tree
- Leaf nodes represent actual rectangles to be indexed
- Internal nodes represent smallest axes-parallel rectangle containing all descendents
- Rectangles at any level may overlap
- Good subdivisions:
 - Minimize the total area of containing rectangles
 - Minimize the total area of overlap of containing rectangles
- Overlap is critical: point and range searches are inefficient with large overlap (R+-tree aims to eliminate overlaps)

R-Tree



Ordered and Unordered Files

- ◆ In **unordered files** new records are inserted in the next physical location on the disk
 - ◆ Insertion is very efficient
 - ◆ Retrievals require search through every record in sequence: **linear search** with time complexity $O(n)$
 - ◆ Deletion causes “holes” to appear in sequence
- ◆ In **ordered files** each record is inserted in the order of the values of one or more of its fields
 - ◆ Slows the insertion of new records
 - ◆ Allows efficient **binary search** with time complexity $O(\log_2 n)$ on indexed field, but not on other fields

Two- Dimensional Ordering

From one to two dimensions


The main problem facing multidimensional spatial data structures is that data storage is essentially one-dimensional

- ◆ Many common indexes assume a grid-based representation (**tile indexes**)
- ◆ Tile indexes aim to provide a path through the grid that visits each cell
- ◆ Indexes differ in how well they preserve proximity, i.e., cells that are spatially close are close in the index

Summary

- ◆ Physical file organization affects database performance
- ◆ Indexes are needed to go beyond the limitations of physical file organization
- ◆ Non-spatial indexes, like B-trees, are inadequate for storing spatial data
- ◆ The key issue in spatial indexes is representing two dimensional data in a one-dimensional index

Point Quadtree



Spatial Queries


- ◆ **Point query**: retrieve all records with spatial references located at a particular point
- ◆ **Range query**: retrieve all records with spatial references located within a given range (spatial ranges may be any shape, but are often rectangular)

Example

- ◆ Non spatial query: Retrieve the point location of Trentham Gardens
- ◆ Spatial point query: Retrieve any site at location (37, 43)
- ◆ Spatial range query: Retrieve any site in the rectangle defined by (20, 20)–(40, 50)

Region Quadrees

- ◆ Quadrees take full advantage of the spatial structure, adapt to variable spatial detail
- ◆ Inefficient for highly inhomogeneous rasters
- ◆ Very sensitive to changes in the embedding space (e.g., translation, rotation)



Point Quadtree

- ◆ Combination of grid approach with multidimensional binary search tree
- ◆ Each non-leaf node has four descendents
- ◆ Each quadrant partition is centered on a data point
- ◆ Quadtree build time is $O(n \log n)$; search time is $O(\log n)$

